

A Collaborative Filtering Recommender System for Test Case Prioritization in Web Applications

Abstract—The use of relevant metrics of software systems could improve various software engineering tasks, but identifying relationships among metrics is not simple and can be very time consuming. Recommender systems can help with this decision-making process; many applications have utilized these systems to improve the performance of their applications. To investigate the potential benefits of recommender systems in regression testing, we implemented an item-based collaborative filtering recommender system that uses user interaction data and application change history information to develop a test case prioritization technique. To evaluate our approach, we performed an empirical study using three applications with multiple versions and compared five control techniques. Our results indicate that our recommender system can help improve the effectiveness of test prioritization; the performance of our approach was particularly noteworthy when we were under a time constraint.

Keywords: Recommender system, test case prioritization, regression testing, risk measurement, code quality.

I. INTRODUCTION

Software systems undergo many changes during their lifetime, and often such changes can adversely affect the software. To avoid undesirable changes or unexpected bugs, software engineers need to test the overall functionality of the system before deploying a new release of the product. One of the common ways to evaluate system quality in a sequence of releases is regression testing. In regression testing, software engineers validate the software system to ensure that new changes have not introduced new faults or that they don't affect the other parts of the system. However, modern software systems evolve frequently, and their size and complexity grow quickly, and thus the cost of regression testing can become too expensive [4]. To reduce regression testing cost, many regression testing and maintenance approaches including test selection and test prioritization [40] have been proposed.

To date, most regression testing techniques have utilized various software metrics that are available from software repositories, such as the size and complexity of the application, code coverage, fault history information, and dependency relations among components. Further, various empirical studies have shown that the use of a particular metric or combination of multiple metrics can improve the effectiveness of regression testing techniques better than others. For example, Anderson et al. [5] empirically investigated the use of various code features mined from a large software repository, showing that these code features can improve regression testing processes. However, we believe that, rather than simply picking one metric over another, adopting a recommender system, that identifies more relevant metrics by considering software characteristics

and the software testing environment might provide a better solution.

Recommender systems have been utilized to alleviate the decision making effort by providing a list of relevant items to users based on a user's preference or item attributes. For example, companies that produce daily-life applications, such as Netflix, Amazon, and many social networking applications [19], are adopting recommender systems to provide more personalized services so that they can attract more users. Recently, recommender systems have been used in software engineering areas to improve various software engineering tasks. For example, Anvik et al. conducted research that applied machine learning techniques to developers as well as bug history to make suggestions about "who should fix this bug?" [6]. While many software engineering techniques have started to incorporate recommendation systems, no researchers have investigated the use of recommender systems in the area of regression testing.

Therefore, we have investigated whether the use of recommender systems can improve regression testing techniques, in particular focusing on test case prioritization. To implement the recommender system, we used user interaction data and application change history information. Previous studies have shown that change history information is an effective indicator for bug prediction [28], [21], and the most frequently accessed components have a higher impact on the user-perceived reliability of the application [35], [4]. Using this kind of information, our recommender system identifies potential components that contain faults that can lead to system failure. By running test cases that exercise such components earlier, we can in turn find defects that are exposable by user interactions earlier. We implemented a test case prioritization technique by applying our recommender system and performed an empirical study using two open source and one commercial software system. The results of our study show that our proposed recommender system approach can improve the effectiveness of test case prioritization compared to five other control techniques. Further, our results indicate that our proposed approach was more effective than traditional prioritization techniques when companies have limited time budgets.

The rest of the paper is organized as follows. In Section II, we discuss the approach used in this research and formally define collaborative filtering recommender systems. Sections III and IV present our empirical study, including the design, results, and analysis. Section V discusses the results and the implications of these results. Section VI discusses potential threats to validity, and Section VII presents background and

related work. Finally, in Section VIII, we provide conclusions and discuss future work.

II. THE PROPOSED APPROACH

In this section, we describe the proposed approach, which utilizes an item-based recommender system to improve test case prioritization techniques. Figure 1 shows three major activities in our approach and how these activities are related to each other. The first step of our proposed technique is usage pattern extraction, which is shown in the upper box of the Figure 1. In this activity, we analyze the users' interaction data to determine the most frequently accessed components of the system. Our second activity is change history analysis, which is shown in the lower box of the Figure 1. In this activity, we build a classification model to measure the relationship between software defects and change history metrics. Then, by obtaining the output from these two activities, we measure the risk score of each component and prioritize the test cases based on their coverage of these risky components. In this paper, component refers to method.

A. Usage Pattern Extraction

The goal of our recommender system is to suggest the highest-risk components with the most access frequency among all other components in the applications. In large scale applications, there are wide range of features and components; however, in reality, only a relatively small subset of components is accessed by users. Therefore, even if there are bugs in some part of the system that is not generally access by users, we assume that those bugs have less impact on the user-reliability perception of a system; catching those bugs that have been exposed by users is a bigger priority. Also, prioritizing the test cases that exercise these more frequently accessed components can speed up the rate of fault detection.

There are two primary techniques for collaborative filtering algorithms: user-based and item-based algorithms. A user-based collaboration filtering algorithm tries to find users who are most similar to the current active user. In an item-based algorithm, first we design a model of user rating, and then we evaluate the expected ratings of an item based on the previous ratings of the other similar items. Because our concern is to

find the access frequency of the components rather than user preferences, we chose an item-based algorithm.

To calculate the access frequency of the components, we used two collected data sets: a list of users $U = \{u_1, u_2, \dots, u_m\}$ and a list of the components $C = \{c_1, c_2, \dots, c_n\}$. Each component c_i has a list of users' ratings if users have performed at least one task with it. Typically, in recommender systems, prediction is based on the numerical values of ratings from active users, but in our case we do not have access to such rating modules; instead, we refer the value of access frequency for a specific component by an individual active user as a rating score.

Suppose that a user has interacted with a set of components, $\{c_1, c_2, \dots, c_n\}$. The item-based collaborative algorithm computes the similarity between components by calculating the cosine similarity of the ratings of two components. Then, the algorithm selects the N most similar components that has the highest similarity score to c_i .

After finding the components that have the highest similarity score, we need to predict the ratings of those components that have been ignored or have not yet been used. To do this, the algorithm uses a weighted average of the target user ratings on these similar components. Below, we describe the process of computing similarity and ranking.

Suppose that we have a web application that has several functionalities; a group of users shows similar interests in using a set of components, while other groups of users use different sets of components. We want to measure the similarity between components by considering users' activities and their preferences in using the components.

Figure 2 illustrates an example of component similarity prediction. The upper left-hand matrix in the figure shows how many times each component has been used by users. The numbers in the cells show the access frequency by the user u_i on the component c_j , and NA indicates that the user u_i has not used that particular component yet. In this figure, u_1, u_3, u_6 , and u_8 utilized components c_1 ; and u_1, u_2, u_4, u_5 , and u_9 utilized c_2 . We refer to a set of component ratings as a component vector. For example, the vector c_1 is $\{5, NA, 6, NA, NA, 3, NA, 1, NA, NA\}$. Once we identify the component vectors, we measure the similarity between c_i and all other components one by one.

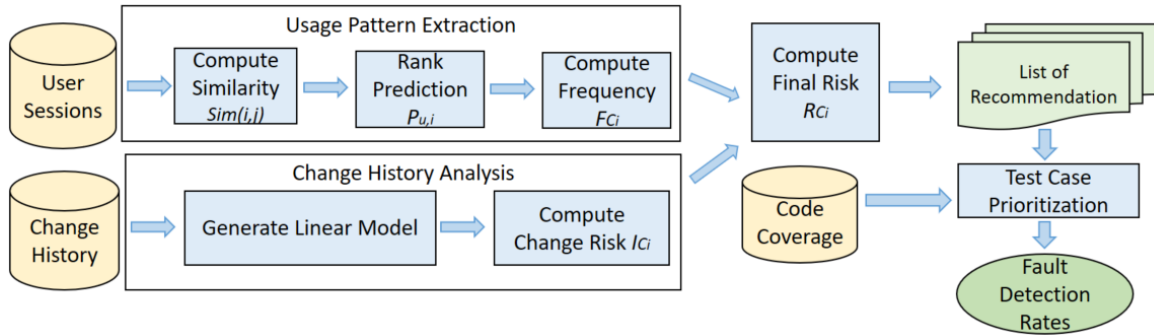


Fig. 1: An Overview of Our Test Case Recommender System

In order to determine the similarity between two components i and j , we use Pearson-r correlation [36]. If U is the set of users who rated components i and j , then we compute the correlation similarity using the following equation:

$$Sim(i, j) = \frac{\sum_{u \in U} (C_{u,i} - \bar{C}_i)(C_{u,j} - \bar{C}_j)}{\sqrt{\sum_{u \in U} (C_{u,i} - \bar{C}_i)^2} \sqrt{\sum_{u \in U} (C_{u,j} - \bar{C}_j)^2}}$$

where $C_{u,i}$ is the value of access frequency for component i by user u , and \bar{C}_i is the average access frequency value of C_i .

The reason for finding the similarity between components is to find the missing value in the matrix. To measure the similarity between components, we normalize the rating of each component by subtracting the row mean. The upper right-hand matrix in Figure 2 shows the normalized rating values of the left matrix. For example, the average of access frequency of c_1 and c_2 is $15/4$ and $26/5$, respectively. Then, we subtracted mean values from each corresponding row of the left matrix. Positive values indicate that the user liked the component more than average; while negative values indicate that the user liked the component less than average; 0 indicates the average access frequency for a component. We treat blank values as 0. The rightmost column in the upper right-hand matrix shows the similarity between c_1 and all other components. For example, $Sim(c_1, c_2) = -0.54 < Sim(c_1, c_4) = 0.38$ means that the probability of rating of c_1 is much more like c_4 than c_2 .

After calculating the similarity between all components, we select the set of N most similar components for c_i ; this process will be iterated for every component. Once we have this set of N , then we can make a prediction of access frequency for the missing values of c_i based on a rating of the N similar components. To estimate the access frequency rates of ignored components, we performed ratio prediction computation using a weighted sum method. This method provides the ratio

prediction of a specific component i for user u based on similar components.

$$P_{u,i} = \frac{\sum_{all\ similar\ components, N} (S_{i,N} * R_{u,N})}{\sum_{all\ similar\ components, N} (S_{i,N})}$$

$R_{u,N}$ is the rating score for user u and N most similar components, and $S_{i,N}$ is the similarity score of component i and N .

We illustrate this using an example in Figure 2. The lower matrix in the figure shows the predicted frequency access for $P_{1,4}$, which is calculated by $(0.38*6/5 + 0.18*-13/7)/(0.38 + 0.18) = 0.23$.

The number, N , is determined based on the context of application domain, ranging from 1 to size of $C - 1$. However, assigning a large number to N will increase the calculation cost significantly, while the result accuracy would not change noticeably. Therefore, to reduce the cost of the calculation process, we set $N = 2$, which means that to predict the missing values we only select the two most similar components to c_i . This process is repeated until we find the ranking for the all missing values. Once we calculate ratio scores for the components, we can produce a matrix of components and their access frequency ranking scores. We calculate normalized access frequency scores using the following equation:

$$F_{Ci} = \frac{\sum (P_{Ci})}{number\ of\ components}$$

where P_{Ci} is the predicted rank score and F_{Ci} is the normalized rank score of component i . Then we can sort the matrix of components based on their ranking to select $Top-N$ most frequently accessed components.

B. Change History Analysis

The second phase of our proposed approach is change history analysis. Among hundreds of attributes of code and change history metrics to evaluate the code quality and error

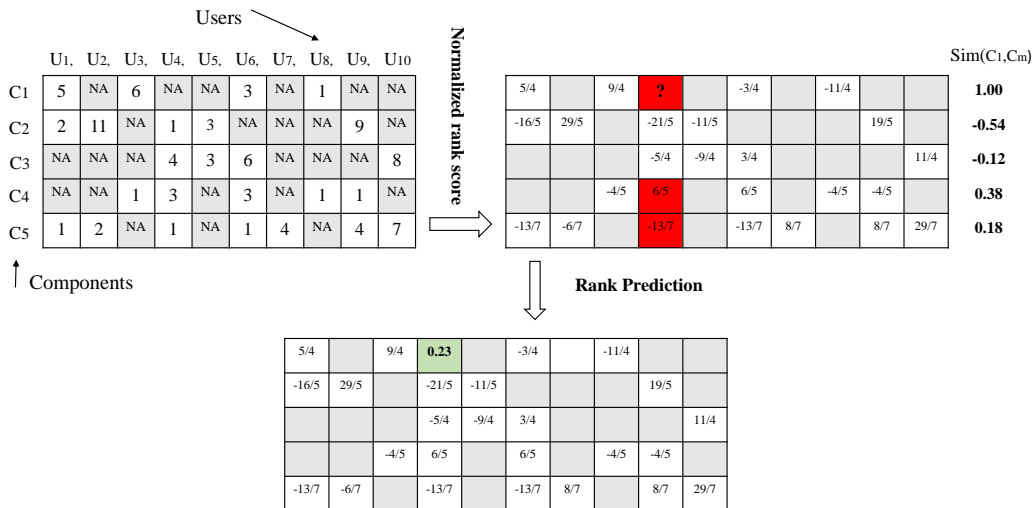


Fig. 2: Item-Based Collaborative Filtering Process for Identifying Most-Frequent Components

proneess, we chose change history metrics to identify the riskiest components. According to a previous study [28], change history is a better indicator for bug prediction purposes than code metrics.

The process of change analysis involves two major steps. First, we collect change history information (e.g., added lines of code, deleted lines of code, etc.) and bug reports from all available versions of the applications from their repositories. The details of collecting the change history information are discussed in Section III-C. Once the change history data is collected, we then design a linear model from a set of collected change metrics to build a classification model for bug prediction. Our goal is to find the correlation coefficient of each metric to measure statistical relationships between a change metric and real defects. The value of this measure ranges between 1 and -1, where 1 indicates a strong positive relationship, 0 indicates no correlation, and a negative value indicates a reverse correlation.

To evaluate our linear model, we applied 10-fold cross validation and repeated this process 100 times. We used a common accuracy indicator to determine the accuracy of our model. The three accuracy indicators that we used are PC, TP and FP. PC indicates the percentage of correctly classified instances, TP (true positive) indicates the number of components that contain a bug (and our classification model also classified them as buggy components), and FP (false positive) is the number of components that are classified as buggy (but they are clean). The output of our classification model is a list of components with their risk values, I_{Ci} (the risk score of being buggy for component i).

C. Test Prioritization Using the Recommender System

After obtaining the two metrics explained in previous subsections (component risk scores and access frequency ratios), we calculate the final risk scores using the following equation:

$$R_{Ci} = F_{Ci} * I_{Ci}$$

where F_{Ci} is the access frequency score of component i and I_{Ci} is the fault risk score of C_i .

Using R_{Ci} scores, our recommender system provides a ranked list of components. The ranked list of components contains those components of the system that are most likely to be the cause of regression faults. As shown in Figure 1, the test case prioritization algorithm reads two inputs (a recommended $Top - N$ list of components and code coverage of tests), and reorders test cases.

For example, suppose we have five components $C = \{c_1, c_2, \dots, c_5\}$ with risk scores $R = \{0.0014, 0.251, 0.034, 0.561, 0.138\}$. Also, suppose we have a list of test cases with their code coverage information, such as $TC = \{T_1 = (c_5, c_3), T_2 = c_1, T_3 = (c_4, c_2), T_4 = (c_2, c_5), T_5 = c_2\}$. Then, we reorder the test cases in this order: T_3, T_4, T_1, T_2, T_5 , since T_3 covers c_4, c_2 , the two components having the highest risk score ($0.561 + 0.251 = 0.812$), and T_2 and T_5 will be executed last because T_2 covers c_1 , which has the lowest risk score (0.0014). T_5 only

TABLE I: Application Properties

Metrics	DASCP	nopCommerce	Coevery
Classes	107	1,919	2,258
Files	201	1,473	1,875
Functions	940	21,057	13,041
LOC	35,122	226,354	120,897
Sessions	748	1310	274
Faults	35	70	30
Versions	3	23	3
Test Cases	95	543	1,120
Installations	3	2	1

covers c_4 (this component has been covered by T_3 as well); and because T_4 covers more components ($0.251 + 0.138 = 0.389$), then it has higher priority than T_5 , which covers only $c_2 = 0.251$. Finally, we calculate the fault detection rate of the reordered test cases by applying the proposed technique in the latest version of each application.

III. STUDY

Our study investigates whether the use of a recommender system can improve the effectiveness of test case prioritization techniques considering the following research questions.

- RQ1: Can our recommender system improve the effectiveness of test case prioritization techniques?
- RQ2: Can our recommender system be effective in improving the effectiveness of test case prioritization techniques when we have a limited time budget?

The following subsections present our objects of analysis, study setup, threats to validity, and data analysis.

A. Objects of Analysis

To investigate our research questions, we performed an empirical study using two open source applications and one commercial web application.

DASCP is a digital archiving and scanning software designed for civil projects; we obtained this application from a private company. DASCP is a web based application written in ASP.Net and designed to store civil project contracts, which include the technical information of civil and construction projects such as project plans and relevant associated information. Our second application is **nopCommerce**, which is a widely used open source e-commerce web application with more than 1.8 million downloads. This application is written in ASP.Net MVC and uses Microsoft SQL Server [1]. The last application is **Coevery**, which is an open source customer relationship management (CRM) system written in ASP.Net. This application provides an easy framework for users to create their own customized modules without having to write any code [2].

Table I lists the applications under study and their associated data: “Classes” (the number of class files), “Files” (the number of files), “Functions” (the number of functions/methods), “LOC” (the number of lines of code), “Sessions” (the number of user sessions that we collected), “Faults” (the total number of seeded and real faults), “Version” (the number of versions), “Test Cases” (the number of test cases), and “Installations” (the number of different locations where the applications were

installed). Test cases were in the application package; we did not generate any new test cases. We downloaded all available versions of open source applications from the applications' *GitHub* repositories.

B. Variables and Measures

1) *Independent Variable*: To investigate our research questions, we manipulated one independent variable: prioritization technique. We considered six different test case prioritization techniques, which we classified into two groups: control and heuristic. Table II summarizes these groups and techniques. The second column shows prioritization techniques for each group, and the third column is a short description of prioritization techniques.

As shown in Table II, we considered five control techniques and one heuristic technique. For our heuristic technique, we used the approach explained in Section II, so, here, we only explain the control techniques we applied as follows:

Change History-Based (T_{ch}): The first control technique is test case prioritization based on change risk score. In order to perform this technique, we used the information that we obtained from the change history analysis approach, which we explained in Section II-B. We prioritized our test cases based on the highest scores of the change risk list.

Most Frequent Web Forms-Based (T_{mfw}): This approach determines the web forms that have been most frequently used by users. We assume that the most frequently used web forms play a more important role in the system; therefore, they should have a higher priority to be tested first. Another reason why the access frequency of web forms is the key for testing is that the most frequently used web forms usually contain more functionality and links with other forms, so any bugs on those forms can affect a greater portion of the entire system. Here, we define "base session" as a long session performed by our users that displays the most interactions in the application [4]. We picked 20% of our total sessions as base sessions. For example, in online shopping, our base session is a sequence of actions from user login to checkout, including all mandatory actions (e.g., payment) and some non mandatory actions such as browsing for other items, and checking the inbox during the shopping process.

Base Session	Test Sessions
1: PublicStore: Init	2: PublicStore: Init
PublicStore:ViewCategory	PublicStore:ViewCategory
PublicStore:Grid:ViewDetails	DescriptionDialog:OkButton:Click
PublicStore:AddToCompareList	
PublicStore:ViewCategory	3: PublicStore: Init
PublicStore:AddProductToReview	PublicStore:ViewCategory
PublicStore:AddProductToWishlist	PublicStore:AddToCard
PublicStore:AddToCard	Payment:Init
Payment:Init	Payment:ItemsTable:SelectRow
Payment:ItemsTable:SelectRow	DescriptionDialog:OkButton:Click
DescriptionDialog:OkButton:Click	

Fig. 3: Sample of Base Session and Test Sessions

After collecting all sessions, we conducted an analysis of web forms' access frequency by comparing the observed web forms in each session with a base session.

$$F_{w,i} = \frac{\sum_{form\ score\ for\ each\ session}(S_i)}{\sum_{number\ of\ base\ sessions}(BS_i)}$$

Using Figure 3 as an example, the form "PublicStore" from the base session was observed in test sessions 2 and 3. If we assume that we have ten test sessions and three base sessions, and this particular form was observed in seven of the ten forms, then the access frequency of form "PublicStore" is equal to $(0.7/3) = 0.23$.

Most Frequent Methods-Based (T_{mfm}): The most frequent methods approach is nearly identical to the web form frequency technique. The only difference is that we considered a method instead of a web form as a comparison factor. The most frequent methods usually have high dependency on other classes and methods. If one of them fails, it can cause a significant failure or degradation of the system. In order to prevent a domino effect in the system, high-frequency methods should be tested first, because their failure can cause other components to fail due to their dependencies.

Random (T_r): The random prioritization technique randomly reorders test cases.

Greedy (T_g): The greedy technique reorders test cases based on the total number of blocks covered by test cases.

2) *Dependent Variable*: Our dependent variable RQ1 is the average percentage of fault detection (APFD) referring to the average percentage of faults detected during the test suite execution. The range of APFD is from 0 to 100, the higher value indicating better prioritization technique. Given T as a test suite with n test cases and m number of faults, F as a collection of detected faults by T , and TF_i as the first test case that catches the fault i , we calculate APFD [33] as follows:

$$APFD = 1 - \frac{TF_1 + TF_2 + \dots + TF_m}{nm} + \frac{1}{2n}$$

RQ2 seeks to measure the effectiveness of our proposed approach when we have constrained resources, which means that we need to evaluate the effectiveness of our approach using a different metric. Qu et al. [31] defined the normalized metric of APFD, which is the area under the curve when the numbers of test cases or faults are not consistent. The NAPFD formula is as follows:

$$NAPFD = p - \frac{TF_1 + TF_2 + \dots + TF_m}{nm} + \frac{p}{2n}$$

In this formula, n is a percentage of the test suite run, m represents the total number of faults found by all test cases, TF_i indicates the same parameter as APFD, and p is the number of faults detected by the percentage of our budget divided by the total number of detected faults when running 100% of test cases.

TABLE II: Test Case Prioritization Techniques

Group	Technique	Description
Control	Change history-based (T_{ch})	Test case prioritization based on change risk analysis score.
	Most frequent web forms-based (T_{mf_w})	Test case prioritization based on value of access frequency for each web form.
	Most frequent methods-based (T_{mf_m})	Test case prioritization based on value of access frequency for each method.
	Random (T_r)	Test case prioritization in random order.
	Greedy (T_g)	Test case prioritization based on the code coverage.
Heuristic	Hybrid collaborative filtering-based (T_{hcf})	Test case prioritization based on the proposed technique.

C. Data Collection and Experimental Setup

In order to perform our experiment, for both the control and heuristic techniques we needed to collect three different types of datasets: telemetry data, change history, and code coverage information. We explain the data collection processes in the following subsections.

1) *Collection of Telemetry Data*: To collect telemetry data, we implemented a small function to record user interactions. We considered a sequence of each user's interactions on a specific date as a user session. First, we uploaded two applications, *Covery* and *nopCommerce*, on an IIS server at the University of ABC in November 2016. The server specification is CPU Core i7, with 16 GB of RAM. After deploying our applications, we recruited volunteer graduate and undergraduate computer science students and assigned a variety of tasks to them. These tasks to the volunteers were simple scenarios that each application is designed for. For example, in *nopCommerce*, we asked the volunteers to perform online shopping following the actual necessary steps, from login to payment. We also asked some of the users to be the system administrator, so we could monitor the whole system rather than only the end user side. We asked the end users to access other part of the system randomly, for example, checking their inbox or wish list. In total, seventy volunteer students performed different tasks during a forty-day period. We collected 1,310 and 274 user sessions for *nopCommerce*, and *Covery* respectively.

The data collection process for *DASCP* is different from that of *Covery* and *nopCommerce*. *DASCP* is a commercial and closed source application that has three versions, and it has been installed on the servers of three companies since 2011. The *DASCP* users whose data we examined are real users, and they have application domain knowledge. We collected a twelve-month period of user interactions for *DASCP*. As described in Section III-A, *DASCP* provides two types of access rights for users. In this study, we only considered the sessions of those users who have full access to the system. In total, 748 user sessions were collected during the designated period of time.

However, the length of the sessions varied by user, date, and workplace. For example, some users performed all their assigned task a few hours before the determined deadline, while others distributed their tasks over several days. The average length of user sessions for *nopCommerce* is equal to 56 and, for *Covery*, 24. However, in some case we obtained session lengths over 300, especially when the interaction date was close to the deadline. Also, most of the *Covery* users were selected among graduate students, since the functionality

of this application is relatively more complex than of *nopCommerce*.

Figure 4 shows an example of the raw telemetry data. The left-hand column shows the session identifier, which is a user navigating through the system. The right-hand column is the set of server-side user interactions. The structure of the interactions is of the format (Form name):(Control name):(Action name).

0E98L725-M51C-4BFA-9960-E1C80M27ABA0	PublicStore:PublicStore:Init
0E98L725-M51C-4BFA-9960-E1C80M27ABA0	PublicStore:ViewCategory:Init
0E98L725-M51C-4BFA-9960-E1C80M27ABA0	PublicStore:ItemsGrid:ViewDetails
0E98L725-M51C-4BFA-9960-E1C80M27ABA0	PublicStore:AddProductToReview:Click
0E98L725-M51C-4BFA-9960-E1C80M27ABA0	PublicStore:AddProductToWishlist:Click
0E98L725-M51C-4BFA-9960-E1C80M27ABA0	PublicStore:AddToCard:Click
0E98L725-M51C-4BFA-9960-E1C80M27ABA0	PublicStore:Payment:Init
0E98L725-M51C-4BFA-9960-E1C80M27ABA0	Payment:ItemsTable:SelectRow
0E98L725-M51C-4BFA-9960-E1C80M27ABA0	Payment:DescriptionDialog:Click
0E98L725-M51C-4BFA-9960-E1C80M27ABA0	Payment:OkButton:Click

Fig. 4: Sample User Session

2) *Collection of Code Change History*: We had to take three steps to measure change risk. First, we needed a clear understanding of the applications with respect to their changes. For instance, we needed to check whether a change was just the renaming a variable or component, the addition of some comments, or an alteration of code by adding or deleting functions, and so on. Then, we needed to check whether changes had been made in the current version, and finally, we tested a recently changed system [7]. Figure 5 shows the versions that we used in this study. *nopCommerce* contains 36 versions but we used only the versions available on the *GitHub* repository as of the experiment date. However, for *Covery* and *DASCP*, we used all available versions. All studied applications contain fine-grained changes, and the commits on these two open source applications are available in the *GitHub* repository.

In our study, we collected the change history of our three applications following Giger et al.'s approach [21]. Most of these metrics have been used in bug detection research, and they are known to be good indicators for locating bugs [26], [38], [28], [24], [30], [21]. Table III shows the applied change metrics in this study.

3) *Collect Code Coverage*: Once our recommender system was designed and implemented, we needed to find test cases that covered the recommended components. We collected the code coverage data for our test cases using the code coverage analysis tool that Microsoft Visual Studio provides as part of its framework. After collecting the code coverage information,

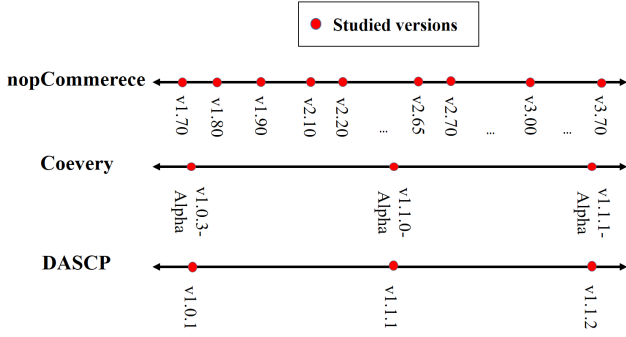


Fig. 5: Versions of Each Application with Change Information and Bug Reports

TABLE III: Change Metrics Used to Evaluate Risk

Metrics Name	Description
Modification	Number of times a method was changed
LOC Added	Added lines of code to a method body over all method histories
Max LOC Added	Maximum added lines of code to a method body for all method histories
AVE LOC Added	Average added lines of code to a method body per method history
LOC Deleted	Deleted lines of code from a method body over all method histories
Max LOC Deleted	Maximum deleted lines from a method body for all method histories
AVE LOC Deleted	Average deleted lines from a method body per method history
Code Churn	Sum of all changes over all method histories
Max Code Churn	Maximum code churn for all method histories
AVE Code Churn	Average code churn per method history
Age	Age of a method in weeks from last release

we entered it into a relational database. We assigned unique identifier values for each method and test case, which provides a key for method and test tables. Hence, we could easily map the methods to the test cases that exercise them.

TABLE IV: Code Coverage Data Table

MethodID	Risk Score	TestID1	...	TestID n
12	0.876	0		0
287	0.012	1		0
301	0.547	0		0
148	0.145	0		1
67	0.055	1		0

Table IV show the code coverage data we collected. The first column, “MethodID”, shows the unique identifier values that we assigned for each method. The second column shows the final risk scores, which is the output of our recommender system. Other columns list our test cases with Boolean values: 0 indicates that the test case does not cover the method, and 1 indicates that the test case covers the method.

4) *Faults*: The applications contain real defects reported by users, but the number of defects is rather small considering the size of the applications we studied. Thus, two graduate students seeded additional faults by simulating programmers’ common mistakes (e.g., using a single equal sign to check equality). All seeded faults were in the server-side code level, and we ignored HTML-based and GUI faults. Four types of faults were seeded: (1) data faults that are related to interaction

with the data store; (2) logic faults that are logic errors in the code; (3) action faults that modify parameter values and actions; and (4) linkage faults that change the hyperlinks references. Table I shows the total number of faults that including both real and seeded faults. The number of real faults for *nopCommerce*, *Coevery*, and *DASC* is 25, 8, and 13, respectively. Once we collected all the required data, we ran control and heuristic techniques and calculated all the dependent values for the reordered test cases.

IV. DATA AND ANALYSIS

In this section, we present the results of this study, considering each of the two research questions.

A. RQ1 Analysis

RQ1 investigates whether the use of the recommender system can improve the effectiveness of test case prioritization techniques. Figure 6 presents the results for *DASC*. The vertical axis shows APFD scores, and the horizontal axis shows prioritization techniques: “ T_{ch} ” (change history-based), “ T_{mfw} ” (frequent web form-based), “ T_{mfm} ” (frequent method-based), “ T_r ” (random order), “ T_g ” (greedy technique), and “ T_{hef} ” (recommender system-based).




Fig. 6: DASC: APFD Boxplots

As can be seen the heuristic (T_{hef}) outperformed all control techniques, but data distribution patterns are different. Among control techniques, the random technique was the worst performer, whereas the frequent method-based technique (T_{mfm}) produced results similar to the change history based technique (T_{ch}).

Figure 7 presents the results for *nopCommerce*. The results show that our proposed technique outperformed all five control techniques, but the variance among the APFD scores is relatively high. Among the first three control techniques, the change history-based technique (T_{ch}) produced slightly better results than the others. This application has 23 versions, and we had a sufficient amount of change history data for our training, so we speculate that this is the reason why the change history-based technique produced better results than the other control techniques. The greedy technique, which is the fifth control technique, did not yield better results than the other control techniques. The poor performance relative to other techniques could be explained by the nature of test cases of this application. A large number of test cases have

high code coverage, which caused the greedy technique which prioritizes based on block code coverage, to perform almost like the random technique.




Fig. 7: nopCommerce: APFD Boxplots

Figure 8 shows the results for *Coevery*. When we compared the median values, our heuristic technique improved the rate of fault detection by 21% over the four control techniques (T_{ch} , T_{mfw} , T_{mfm} , and T_g) on average and 41% over the random technique. Among the first three control techniques, T_{ch} and T_{mfw} produced slightly better results than T_{mcf} .




Fig. 8: Coevery: APFD Boxplots

Overall, the results showed that our proposed recommender system-based technique performed better than all the control techniques across all three applications. Among the control techniques, the random technique performed worst across all applications, whereas the change history-based technique produced slightly better and more stable results for all applications. Also, in the case of *DASCP* and *Coevery*, the greedy technique (T_g) showed marginally better results than other control techniques.

B. RQ2 Analysis

RQ2 investigates whether the use of the recommender system can improve the effectiveness of test prioritization when we have a limited budget for testing, a common situation that the software industry often faces.

In RQ2 analysis, we measured the NAPFD, which is a normalized ratio of APFD, when our resources were not consistent. In this experiment, first we executed 10% of our test cases, and we continued to execute the test cases in increments of 10% of the total until they had all been executed to see

whether we could improve the fault detection rate, given a time constraint dictating that running 100% of the test cases at one time was not feasible.

Table V shows the results of our three applications. This table shows the results of two primary analyses. The first part of the table presents the NAPFD scores on average, and the second part shows the improvement rates of the heuristic technique over the five other control techniques.

By examining the numbers in the table, we can observe that the improvement rates of our heuristic technique over the control techniques vary widely. When we compared the heuristic with T_{ch} , the improvement rates ranged from 4% to 41% for *DASCP*, from 17% to 63% for *nopCommerce*, and from 17% to 54% for *Coevery*. When compared with T_{mfw} , the improvement rates ranged from 11% to 132% for *DASCP*, from 42% to 140% for *nopCommerce*, and from 14% to 77% for *Coevery*, indicating significant improvements. When compared with T_{mfm} , the improvement rates ranged from 15% to 78% for *DASCP*, from 27% to 78% for *nopCommerce*, and from 27% to 48% for *Coevery*, indicating results similar to those for as T_{ch} . When compared with T_g , the improvement rates ranged from 12% to 41% for *DASCP*, from 24% to 129% for *nopCommerce*, and from 13% to 67% for *Coevery*, showing improvement over a popular and commonly used technique. As for the comparison with T_r , the results were more remarkable. The rates ranged from 32% to 282% for all three applications.

One outstanding trend we observed in the table is that the improvement rates are much higher when the time budget is smaller. For example, in the comparison with T_{ch} for *nopCommerce*, when 10% of the budget was assigned, the improvement rate was 60%, but when we had a full budget, the rate dropped to 18%. A similar trend can be observed across all control techniques and applications. This indicates that our approach can be more helpful when companies are operating under a tight budget.

V. DISCUSSION

In this section, we discuss the implications of the results for researchers and practitioners as well as the limitations of recommender systems that we found during our experiment.

Implications of the Results: Our experiment results indicate that by utilizing a recommender system when we prioritize test cases, we were able to improve the effectiveness of test case prioritization. Further, we found that our proposed approach is far more effective than the control techniques when testers have a limited time budget during testing. In particular, from the RQ2 results, we learned that our approach produced much greater benefits when the time budget is limited. The findings from this study could help practitioners select appropriate test prioritization techniques considering their product delivery schedule and other circumstances (e.g., the availability of required software attributes).

While we have not conducted a cost benefit analysis of our approach, which is our future work, we found that the cost of applying our approach is negligible. The biggest

TABLE V: NAPFD Scores on Average.

Application	Test Exe.	Techniques						Improvement Rate over Control					
	Rate (%)	T_{ch}	T_{mfw}	T_{mfm}	T_r	T_g	T_{hcf}	T_{hcf}/T_{ch}	T_{hcf}/T_{mfw}	T_{hcf}/T_{mfm}	T_{hcf}/T_r	T_{hcf}/T_g	
DASCP	10	18.54	12.17	15.12	14.33	16.53	23.37	26%	92%	54%	63%	41%	
	20	21.31	12.88	16.78	15.51	25.13	29.98	40%	132%	78%	93%	19%	
	30	28.86	21.19	25.12	18.47	36.42	40.95	41%	93%	63%	121%	12%	
	40	40.42	29.86	38.68	25.59	41.02	55.3	36%	85%	42%	116%	34%	
	50	54.34	36.21	42.9	39.15	50.13	67.07	23%	85%	56%	71%	33%	
	60	61.7	47.31	50.34	45.29	54.38	70.42	14%	48%	39%	55%	29%	
	70	68.34	56.96	65.97	60.74	62.36	76.67	12%	34%	16%	26%	22%	
	80	75.41	69.04	71.14	64.88	72.4	84.01	11%	21%	18%	29%	16%	
	90	83.35	77.28	77.69	67.11	78.4	89.83	7%	16%	15%	33%	14%	
	100	90.16	84.21	79.22	70.91	82.12	94.14	4%	11%	18%	32%	14%	
nopCommerce	10	17.54	12.17	15.75	8.33	12.25	28.14	60%	131%	78%	237%	129%	
	20	29.28	19.88	28.35	12.51	22.39	47.9	63%	140%	68%	282%	113%	
	30	38.86	27.19	35.45	15.57	40.12	55.28	42%	103%	55%	255%	37%	
	40	44.42	34.86	48.68	27.59	48.66	65.06	46%	86%	33%	135%	33%	
	50	58.34	39.16	54.94	35.91	53.26	74.78	28%	90%	36%	108%	40%	
	60	62.7	51.42	57.08	41.45	60.62	81.49	29%	58%	42%	96%	34%	
	70	68.14	55.03	59.97	50.02	69.23	86.38	26%	56%	44%	72%	24%	
	80	76.22	60.21	64.14	58.15	71.5	89.87	17%	49%	40%	54%	25%	
	90	79.4	66.01	70.22	60.17	73.66	95.14	19%	44%	35%	58%	29%	
	100	82.16	68.32	76.04	63.91	74.13	97.06	18%	42%	27%	51%	30%	
Coevery	10	26.54	23.17	29.12	20.33	24.45	41.02	54%	77%	40%	101%	67%	
	20	44.28	42.88	45.12	24.51	42.12	66.98	51%	56%	48%	173%	59%	
	30	60.86	44.19	50.12	28.57	58.18	73.28	20%	65%	46%	156%	25%	
	40	63.42	51.51	53.68	31.59	63.4	75.06	18%	45%	39%	137%	18%	
	50	64.34	55.74	58.3	39.62	64.73	77.1	19%	38%	32%	94%	19%	
	60	66.7	56.33	60.41	44.09	66.23	78.65	17%	39%	30%	78%	18%	
	70	68.19	59.86	62.97	46.11	69.07	81.13	18%	35%	28%	75%	17%	
	80	70.67	66.1	65.14	57.91	73.19	83.2	17%	25%	27%	43%	13%	
	90	71.81	73.14	66.03	59.01	74.36	86.69	20%	18%	31%	46%	16%	
	100	72.16	76.21	68.22	62.91	75.15	87.23	20%	14%	27%	38%	16%	

cost for applying our approach is calculating the matrix of access frequency scores, but it took only 580, 153, and 374 seconds for *nopCommerce*, *Coevery* and *DASCP*, respectively. Further, the prioritization algorithm took only 0.29 seconds for *nopCommerce*, the application that has the largest number of test cases.

Limitations of Applying Recommender Systems: While our results indicate that recommender systems can improve test case prioritization, like any other approaches, recommender systems also have their own limitations. In this study, we applied an item-based collaborative filtering algorithm that calculates the components' similarity based on user ratings. This means that our results are highly dependent on the accuracy of user ratings, and thus user ratings can greatly affect the experiment's results. There are three common limitations in collaborative filtering recommender systems; new user problem, new item problem, and sparsity [19]. Among these three limitations, two of them are related to our study.

Sparsity. As we discussed earlier, we do not have a rating module in our applications, and thus we used access frequencies to each component as a user rating. Further, for the open source applications, *nopCommerce* and *Coevery*, we collected the user interaction data from nonprofessional users, so our data could contain noise and redundancy, which can affect the performance of our technique. Moreover, the number of collected ratings is relatively small, compared to the number of expected ratings required to generate an accurate prediction.

Also, the distribution patterns of user ratings can affect the outcome of collaborative filtering algorithms. For example, in *nopCommerce* case, some components were used by all

users, such as registration and membership components, while some other components were ignored by the majority of users. In order to eliminate this issue, we need to collect more user interaction data by considering a larger number of users and a longer period of time. In addition, having actual users would generate more realistic results, because their interactions would be based on real business functions and system workflows.

New Item Problem. Another limitation of collaborative filtering that is related to our study is the "New Item Problem". It is a common practice that newly developed components are frequently added to a system. However, rankings on collaborative filtering algorithms are based on user access frequencies to the components. Therefore, the newly added components would not appear in the recommender suggestion list until a certain number of users perform some tasks on them. In our study, in addition to using access frequency scores, we also applied change risk scores to recommend the riskiest components. Therefore, even if a new component has a high change risk score, its frequency score would still be zero, which makes the overall risk score zero. To overcome this limitation, we need to use a hybrid and normalized ranking score by assigning a small value to the components.

VI. THREATS TO VALIDITY

The primary threat to validity of this study is the amount of user session data and the type of users who participated in this study. For the proprietary application that we used in this study, we collected user interaction data for a long period time; the collected data was created by actual users of the application. However, for the two open source applications, the

period of time over which we collected user interactions was relatively short, and the participants were not domain experts or regular users of the applications, so their usage patterns had wide variations. This threat can be addressed by performing additional studies that monitor user interactions over a longer time period among a wider population, and by considering industrial applications and different types of applications (e.g., mobile applications).

Another concern involves the bug reports that we used. Our classification prediction values for designing linear models in the change history analysis were generated from bug history that was reported by actual users. Further, using these bug reports, we measured the coefficient of other variables to create our linear model for change history analysis. Because our bug report data is not comprehensive and contains only those bugs accrued up to the time that we stopped collecting data, and because there might be other bugs that have not been reported yet or that might occur later, there is a possibility that the bug reports are biased.

VII. RELATED WORK

This section discusses two topics related to our research: test case prioritization techniques and recommender systems.

Test Case Prioritization: Due to the appealing benefits of test case prioritization in practice, many researchers have proposed various techniques. These techniques help engineers discover faults early in testing, which allows them to begin debugging earlier. Recent surveys [10], [40] provide a comprehensive understanding of overall trends of the techniques and suggest areas for improvement. Depending on the types of information available, various test case prioritization techniques can be utilized, but the majority of prioritization techniques have used source code information to implement the techniques. For instance, many researchers have utilized code coverage information to implement prioritization techniques [16], [25], [34]. Although this approach is naïve, many empirical studies have shown that this approach can be effective [14], [15], [17], [31].

Further, some researchers have used software risk information in testing approaches. For instance, Frankl and Weyuker [18] introduced two risk related measures of software testing effectiveness, which are expected detected risk and expected risk reduction and investigated the effectiveness of these two measures on testing techniques. Hettiarachchi et al. [22] proposed a new test case prioritization technique. Their technique uses risk levels of potential defects to detect risky requirements then it prioritizes test cases by mapping the related test cases and these requirements.

More recently, several prioritization techniques utilizing other types of information have also been proposed. For example, Anderson et al. applied telemetry data to compute fingerprints to extract usage patterns and for test prioritization [4]. Memon and Amalfitano performed a study applying telemetry data to generate usage pattern profiles [8]. Gethers et al. presented a method that uses textual change of source code to estimate an impact set [20].

Recommender Systems: Recommender systems are software engineering tools that make the decision making process easier by providing a list of relevant items. Some widely-used applications that provide recommender systems are Amazon, Facebook, and Netflix. These applications provide suggestions to target users based on the user or on item characteristic similarities.

With the fast growth of such information, machine learning technologies motivate software engineers to apply recommendation systems in software development. Recommender systems in software engineering have been applied to improve software quality and to address the challenges of the development process [32]. For instance, Murakami et al. [29] proposed a technique that uses user editing activities detecting code relevant to existing methods. Danylenko and Lowe provided a context-aware recommender system to automate a decision-making process for determining the efficiency of non-functional requirements [13].

As discussed briefly earlier, many types of information are available for implementing test case prioritization techniques. In this research, we collected over 2,000 user sessions from three different web applications and gathered the change history of each application. Our research seeks to apply item-based collaborative filtering algorithms to generate a recommendation list for test prioritization. To our knowledge, our recommender system-based prioritization technique is novel and has not yet been explored in regression testing.

VIII. CONCLUSIONS AND FUTURE WORK

In this research, we proposed a new recommender system to improve the effectiveness of test case prioritization. Our recommender system uses three datasets (code coverage, change history, and user sessions) to produce a list of the riskiest components of a system for regression testing. We applied our recommender system using two open source applications and one industrial application to investigate whether our approach can be effective compared to five different control techniques. The results of our study indicate that our recommender system can improve test case prioritization; also, the performance of our approach was particularly noteworthy when we had a limited time budget.

Because our initial attempt to use recommender systems in the area of regression testing showed promising results, we plan to investigate this approach further by considering various algorithms (e.g., a context-aware algorithm and a hybrid algorithm), the characteristics of applications, and the testing context. In future research, we plan to investigate other approaches to address a sparsity problem of our recommender system by applying an associative retrieval framework and related spreading activation algorithms to track user transitive interactions through their previous interactions. To address the “New Item Problem,” we plan to apply knowledge-based techniques such as case-based reasoning. Also, we wish to apply our recommender system to different software domains such as mobile applications, and perform additional studies that monitor user interactions over a longer period of time.

REFERENCES

- [1] <http://www.nopcommerce.com/>. [Accessed: Jan. 26, 2017].
- [2] <http://www.coevery.com/>. [Accessed: Jan. 26, 2017].
- [3] D. Amalfitano, A. R. Fasolino, and P. Tramontana. Rich internet application testing using execution trace data. In *Software Testing, Verification, and Validation Workshops*, pages 274–283. IEEE-ACM, 2010.
- [4] J. Anderson, H. Do, and S. Salem. Customized regression testing using telemetry usage patterns. In *Software Maintenance and Evolution (ICSME), 2016 IEEE International Conference on*. IEEE, 2016.
- [5] J. Anderson, S. Salem, and H. Do. Improving the effectiveness of test suite through mining historical data. In *Working Conference on Mining Software Repositories (MSR)*, 2014.
- [6] J. Anvik, L. Hiew, and G. C. Murphy. Who should fix this bug? In *Proceedings of the 28th international conference on Software engineering*, pages 361–370. IEEE-ACM, 2006.
- [7] Shawn A. Bohner. Extending software change impact analysis into cots components. In *Proceedings of the 27th Annual NASA Goddard Software Engineering Workshop*, page 175. IEEE-ACM, 2002.
- [8] P. A. Brooks and A. M. Memon. Automated gui testing guided by usage profiles. In *Proceedings of the twenty-second IEEE/ACM international conference on Automated software engineering*, pages 333–342. IEEE-ACM, 2007.
- [9] R. Carlson, H. Do, , and A. Denton. A clustering approach to improving test case prioritization: An industrial case study. In *ICSM '11 Proceedings of the 2011 27th IEEE International Conference on Software Maintenance*, pages 382–391. IEEE-ACM, 2011.
- [10] C. Catal and D. Mishra. Test case prioritization: A systematic mapping study. *Software Quality Journal*, 21:445–478, 2013.
- [11] Y. Chen, R. L. Probert, and D. P. Sims. Specification-based regression test selection with risk analysis. In *Proceedings of the conference of the Centre for Advanced Studies on Collaborative research (CASCON)*, 2002.
- [12] K. Christidis, F. Paraskevopoulos, D. Panagiotou, and G. Mentzas. Combining activity metrics and contribution topics for software recommendations. In *Proceedings of the Third International Workshop on Recommendation Systems for Software Engineering*, pages 43–46. IEEE-ACM, 2012.
- [13] A. Danylenko and W. Lowe. Context-aware recommender systems for non-functional requirements. In *Proceedings of the Third International Workshop on Recommendation Systems for Software Engineering*, pages 80–84. IEEE-ACM, 2012.
- [14] H. Do, S. Mirarab, L. Tahvildari, and G. Rothermel. An empirical study of the effect of time constraints on the cost-benefits of regression testing. In *Proceedings of the ACM SIGSOFT Symposium on Foundations of Software Engineering (FSE)*. IEEE-ACM, 2008.
- [15] H. Do, S. Mirarab, L. Tahvildari, and G. Rothermel. The effects of time constraints on test case prioritization: A series of controlled experiments. 36(5), 2010.
- [16] S. Elbaum, A. G. Malishevsky, and G. Rothermel. Test case prioritization: A family of empirical studies. *IEEE Transactions on Software Engineering*, 28(2):159–182, February 2002.
- [17] S. Elbaum, A. G. Malishevsky, , and G. Rothermel. Test case prioritization: A family of empirical studies. 28(2):159–182, 2002.
- [18] P. Frankl and E. Weyuker. Testing software to detect and reduce risk. In *Journal of Systems and Software*, 2000.
- [19] G. and A. Tuzhilin. Toward the next generation of recommender systems a survey of the state of the art and possible extensions. *IEEE Transactions on knowledge and Data Engineering*, 17(6):734–749, 2005.
- [20] M. Gethers, B. Dit, H. Kagdi, and D. Poshyvanyk. Integrated impact analysis for managing software changes. In *Proceedings of the 34th International Conference on Software Engineering*, pages 430–440. IEEE-ACM, 2012.
- [21] E. Giger, M. D’Amorce, M. Pinzger, and H. Gall. Method-level bug prediction. In *ESEM*, 2012.
- [22] C. Hettiarachchi, H. Do, and B. Choi. Effective regression testing using requirements and risks. In *Eighth International Conference on Software Security and Reliability*, pages 157–166. IEEE, 2014.
- [23] D. Jeffrey and N. Gupta. Test case prioritization using relevant slices. In *Int’l Comp. Soft. Appl. Conf.*, pages 411–420, September 2006.
- [24] Y. Jiang, B. Cuki, T. Menzies, and N. Bartlow. Comparing design and code metrics for software quality prediction. In *Proceedings of the 4th international workshop on Predictor models in software engineering*, pages 11–18. ACM, 2008.
- [25] J. Kim and A. Porter. A history-based test prioritization technique for regression testing in resource constrained environments. In *Proceedings of the International Conference on Software Engineering*, May 2002.
- [26] T. Lee, J. Nam, D. Han, S. Kim, and H. Peter. Micro interaction metrics for defect prediction. In *ESEC/FSE ’11 Proceedings of the 19th ACM SIGSOFT symposium and the 13th European conference on Foundations of software engineering*, pages 311–321. IEEE-ACM, 2011.
- [27] S. Mirarab and L. Tahvildari. A prioritization approach for software test cases on Bayesian Networks. In *Found. App. Softw. Eng.*, pages 276–290, March 2007.
- [28] R. Moser, W. Pedrycz, and G. Succì. A comparative analysis of the efficiency of change metrics and static code attributes for defect prediction. In *ICSE ’08 Proceedings of the 30th international conference on Software engineering*, pages 181–190. IEEE-ACM, 2008.
- [29] N. Murakami, H. Masuhara, and T. Aotani. Code recommendation based on a degree-of-interest model. In *Proceedings of the 4th International Workshop on Recommendation Systems for Software Engineering*, pages 28–29. IEEE-ACM, 2014.
- [30] N. Nagappan and T. Ball. Use of relative code churn measures to predict system defect density. In *ICSE ’05 Proceedings of the 27th international conference on Software engineering*, pages 284–292. IEEE-ACM, 2005.
- [31] X. Qu, M.B. Cohen, and G. Rothermel. Configuration-aware regression testing: an empirical study of sampling and prioritization. In *International Symposium on Software Testing and Analysis (ISSTA)*, pages 75–85. IEEE-ACM, 2008.
- [32] M. P. Robillard, W. Maalej, R. J. Walker, and T. Zimmermann. *Recommendation Systems in Software Engineering*. Springer, 2014.
- [33] G. Rothermel, R. Untch, C. Chu, , and M. J. Harrold. Test case prioritization: An empirical study. In *Proceedings of the IEEE International Conference on Software Maintenance*, pages 179–188. IEEE-ACM, 1999.
- [34] G. Rothermel, R. Untch, C. Chu, and M. J. Harrold. Prioritizing test cases for regression testing. *IEEE Transactions on Software Engineering*, 27(10):929–948, October 2001.
- [35] S. Sampath, R. C. Bryce, G. Viswanath, V. Kandimalla, and A. G. Koru. Prioritizing user-session-based test cases for web applications testing. In *ICST ’08 Proceedings of the 2008 International Conference on Software Testing, Verification, and Validation*, pages 141–150. IEEE-ACM, 2008.
- [36] B. Sarwar, G. Karypis, J. Konstan, and J. Riedl. Item-based collaborative filtering recommendation algorithms. In *Proceedings of the 10th international conference on World Wide Web*, pages 285–295. IEEE-ACM, 2001.
- [37] M. Sherriff, M. Lake, and L. Williams. Prioritization of regression tests using singular value decomposition with empirical change records. In *Proceedings of the International Symposium on Software Reliability Engineering*, pages 81–90, November 2007.
- [38] E. Shihab, A. E. Hassan, B. Adams, and Z. M. Jiang. An industrial study on the risk of software changes. In *FSE ’12 Proceedings of the ACM SIGSOFT 20th International Symposium on the Foundations of Software Engineering*, page 62. ACM, 2012.
- [39] R. Robillard Walker and T. Zimmermann. Recommendation systems for software engineering. In *Proceedings of the 27th International Conference on Software Engineering*, pages 80–86. IEEE Press, 2010.
- [40] S. Yoo and M. Harman. Regression testing minimization, selection and prioritization: a survey. *Software Testing, Verification and Reliability*, 22(2):67–120, 2012.
- [41] M. Zanjani, H. Kagdi, and C. Bird. Automatically recommending peer reviewers in modern code review. In *Proceedings of the 37th International Conference on Software Engineering*. IEEE, 2015.